

# Online Learning Applications

## Part 10: Contextual bandits

# Motivating example

Consider a **Movies recommendation** software.

What film should a streaming platform recommend **based on the previous views and ratings?**

# Motivating example

Consider a **Movies recommendation** software.

What film should a streaming platform recommend **based on the previous views and ratings?**

A good system should take into account such information. More in general:

How can we design an algorithm which action (e.g., recommendation) depends on the context (e.g., information about the user)?

# Contextual bandits

## Contextual bandits set-up

For each round  $t \in 1, \dots, T$

- Learner observes **context**  $x_t \in \mathcal{X}$  (**prior to decision!**)
- Learner picks arm  $a_t$
- Learner gets reward  $r_t(a_t) \in [0, 1]$

We will consider the following set-up (many variants are studied in the literature):

- Each  $r_t(a)$  is drawn independently from a fixed **context-dependent** distribution  $\mathbb{P}(\cdot | x_t, a)$ . So  $r_t(a)$  depends both on the context and on the action  $a$
- We denote the expected reward of arm  $a$  under context  $x$  as  $\mu(a|x)$
- The sequence of contexts  $(x_1, \dots, x_T)$  is chosen by an (oblivious) adversary

# Baseline and regret

We consider the following baseline:

## Baseline

The reward of the best policy that maps contexts to actions:

$$\pi^*(x) = \arg \max_{a \in A} \mu(a|x) \quad \forall x \in \mathcal{X}$$

The pseudo-regret with respect to the baseline is:

## Pseudo-regret

$$\mathcal{R}_T = \sum_{t=1}^T (\mu(\pi^*(x_t)|x_t) - \mu(a_t|x_t))$$

Small number of contexts

## Small number of contexts

**Idea:** run a separate copy of a known bandit algorithm (e.g., UCB1) for each context.

---

**Algorithm:** Contextual algorithm

---

```
1 Init: instantiate a (non-contextual) regret minimizer  $\text{ALG}_x$  for each context  $x$ ;  
2 for  $t = 1, \dots, T$  do  
3   | invoke  $\text{ALG}_x$  with  $x = x_t$ , that is: play  $a_t \leftarrow \text{ALG}_x$ ;  
4   | observe  $r_t(a_t)$  and return it to  $\text{ALG}_x$ , that is:  $\text{ALG}_x$  receive as feedback the  
   | reward  $r_t(a_t)$ ;
```

---

# Regret Guarantees

**Assumption:**  $\text{ALG}_x$  has pseudo-regret  $\mathcal{R}_{\text{ALG}_x} \leq O(\sqrt{KT \log T})$ .

## Theorem

The pseudo-regret of the contextual algorithm is

$$\mathcal{R}_T \leq O(\sqrt{KT|\mathcal{X}| \log T}).$$



### Proof sketch.

Let  $n_x = \sum_{t=1}^T \mathbb{I}[x_t = x]$  be the number of times context  $x$  appears. Then, the pseudo-regret accumulated under context  $x$  is at most

$$\mathcal{R}_{\text{ALG}_x} = O(\sqrt{Kn_x \log T}).$$

## Proof sketch.

Let  $n_x = \sum_{t=1}^T \mathbb{I}[x_t = x]$  be the number of times context  $x$  appears. Then, the pseudo-regret accumulated under context  $x$  is at most

$$\mathcal{R}_{\text{ALG}_x} = O(\sqrt{Kn_x \log T}).$$

⚠ The previous bound holds only with high probability (see, e.g., UCB1 proof)

## Proof sketch.

Let  $n_x = \sum_{t=1}^T \mathbb{I}[x_t = x]$  be the number of times context  $x$  appears. Then, the pseudo-regret accumulated under context  $x$  is at most

$$\mathcal{R}_{\text{ALG}_x} = O(\sqrt{Kn_x \log T}).$$

⚠ The previous bound holds only with high probability (see, e.g., UCB1 proof)  
Hence, the overall regret is at most:

$$\sum_{x \in \mathcal{X}} \mathcal{R}_{\text{ALG}_x} = \sum_{x \in \mathcal{X}} O(\sqrt{Kn_x \log T}) \leq O(\sqrt{KT|\mathcal{X}| \log T}),$$

where the last inequality follows from the Cauchy–Schwarz inequality. □

# Drawbacks

Regret bound is very high if  $|\mathcal{X}|$  is large, e.g., if contexts are feature vectors with a large number of features. To handle contextual bandits with a large (or infinite)  $|\mathcal{X}|$ , we either assume some **structure**, or change the **objective**.

Adding structure: Lipshitz contextual bandits

## Adding structure: Lipshitz contextual bandits

Let  $\mathcal{X} \subseteq [0, 1]$  and assume that, for each  $x, x' \in \mathcal{X}$ ,

$$|\mu(a|x) - \mu(a|x')| \leq L |x - x'| \quad \forall a \in \mathcal{A}.$$

### Simple idea:

- Discretize the space of contexts
- Let  $S_\epsilon$  be the  $\epsilon$ -uniform grid on  $[0, 1]$
- Use  $S_\epsilon$  in place of  $\mathcal{X}$
- Map the observed context to the closest point in the grid

We have  $1 + 1/\epsilon$  points in our grid. If this number is small enough we can use the strategy just discussed for small number of contexts.

## Adding structure: Lipshitz contextual bandits

What is the optimal trade-off between:

- having a grid that is precise enough
- having a “small enough” number of points?

We have seen a similar problem regarding pricing.

Setting  $\epsilon = T^{-1/3}$ , we suffer:

- $L \cdot \epsilon \cdot T = L \cdot T^{2/3}$  regret from optimizing over  $L \cdot \epsilon$  optimal solutions
- $O(\sqrt{KT|S_\epsilon| \log T} \simeq T^{2/3})$  regret of the contextual no-regret algorithm

## Adding structure: Lipshitz contextual bandits

What is the optimal trade-off between:

- having a grid that is precise enough
- having a “small enough” number of points?

We have seen a similar problem regarding pricing.

Setting  $\epsilon = T^{-1/3}$ , we suffer:

- $L \cdot \epsilon \cdot T = L \cdot T^{2/3}$  regret from optimizing over  $L \cdot \epsilon$  optimal solutions
- $O(\sqrt{KT|S_\epsilon| \log T} \simeq T^{2/3})$  regret of the contextual no-regret algorithm

This discretization techniques can be applied also to multi-dimensional contexts  $\mathcal{X} \subseteq [0, 1]^d$ , but **the number of contexts increases exponentially in  $d$  providing much worse performances.**



## Adding structure: Lipshitz contextual bandits

What is the optimal trade-off between:

- having a grid that is precise enough
- having a “small enough” number of points?

We have seen a similar problem regarding pricing.

Setting  $\epsilon = T^{-1/3}$ , we suffer:

- $L \cdot \epsilon \cdot T = L \cdot T^{2/3}$  regret from optimizing over  $L \cdot \epsilon$  optimal solutions
- $O(\sqrt{KT|S_\epsilon| \log T} \simeq T^{2/3})$  regret of the contextual no-regret algorithm

This discretization techniques can be applied also to multi-dimensional contexts  $\mathcal{X} \subseteq [0, 1]^d$ , but **the number of contexts increases exponentially in  $d$  providing much worse performances.**

Can we do better in multi-dimensional settings?

## Adding structure: Linear contextual bandits

# Adding structure: Linear contextual bandits

## Contextual linear bandits

For each  $t \in 1, \dots, T$ :

- 1 Observe  $x_t = \{x_{t,a} \in [0, 1]^d \text{ for each } a \in A\}$
- 2 play arm  $a_t \in A$
- 3 get reward  $r_t = \langle \theta^*, x_{t,a_t} \rangle + \epsilon_t$

where

- $\theta^*$  is an **unknown** regression vector
- $\epsilon_t$  is a subgaussian noise

**Goal:** **learn** unknown  $\theta^*$ .... while **maximizing** rewards!

## Lin-UCB (informal)

Linear-contextual bandit problem can be solved by a more complex version of UCB

**idea:** build “confidence region” for the  $\theta^*$  vector

### Idea of Lin-UCB

- In each round we construct a confidence region  $C_t$  such that  $\theta^* \in C_t$  with high probability
- Use  $C_t$  to construct an UCB on the mean reward of each arm given contexts:  
$$UCB_t(a|x_t) = \sup_{\theta \in C_t} \langle x_{t,a}, \theta \rangle$$
- Pick arm  $a$  maximizing  $UCB_t(a|x_t)$

**Known results:** regret UB of  $\tilde{O}(d\sqrt{T})$ , LB of  $\Omega(\sqrt{dT})$  (see, e.g., [Abbasi-Yadkori et al. \[2011\]](#))

## Estimate $\theta^*$

An estimate of  $\theta^*$  can be obtained through least square regression

Design matrix with regularization parameter  $\lambda$ :

$$D_t = \lambda I_d + \sum_{t'=1}^t x_{t',a_{t'}} x_{t',a_{t'}}^\top$$

Regularized least-square estimate:

$$\hat{\theta}_t = D_t^{-1} \sum_{t'=1}^t r_{t'} x_{t',a_{t'}}$$

The estimated reward of an arm  $a$  at round  $t$  is:

$$x_{t,a} = \langle \hat{\theta}_t, x_{t,a} \rangle$$

## Estimate $\theta^*$

An estimate of  $\theta^*$  can be obtained through least square regression

Design matrix with regularization parameter  $\lambda$ :


$$D_t = \lambda I_d + \sum_{t'=1}^t x_{t',a_{t'}} x_{t',a_{t'}}^\top$$

Regularized least-square estimate:

$$\hat{\theta}_t = D_t^{-1} \sum_{t'=1}^t r_{t'} x_{t',a_{t'}}$$


The estimated reward of an arm  $a$  at round  $t$  is:

$$x_{t,a} = \langle \hat{\theta}_t, x_{t,a} \rangle$$

 It is not enough to have an estimate of the rewards. **We need a confidence interval!**

## Confidence sets (informal)

The confidence set  $C_t$  will be an ellipsoid centered in the estimate  $\hat{\theta}_t$ .



$$C_t = \left\{ \theta : \|\theta - \hat{\theta}_t\|_{D_t} \right\} \leq \beta_t$$

Where

- $\|\theta\|_D = \sqrt{\theta^\top A \theta}$
- $\beta_t$  is a parameter

Changing objective: Contextual bandits with a policy class



# Changing the objective: contextual bandits with a policy class

No assumptions on rewards. We make the problem tractable by **restricting the benchmark** in the definition of regret

## Policies and model classes

- A policy  $\pi$  is a mapping from contexts to actions
- A value function  $f$  is a mapping  $f : \mathcal{X} \times \mathcal{A} \rightarrow [0, 1]$  modeling the mean of the rewards distribution (e.g., a linear model, a neural network etc.). We denote by  $\mathcal{F}$  the class of value functions to which we have access
- A value function  $f$  induces a policy  $\pi_f(x) = \arg \max_{a \in \mathcal{A}} f(x, a)$ . The class of policies induced by  $\mathcal{F}$  is  $\Pi$

**Goal:** minimize  $R_T = \sum_{t=1}^T r_t(\pi^*(x_t)) - \sum_{t=1}^T r_t(a_t)$ .

# Connection with supervised ML

Clear connection to “traditional” supervised machine learning: re-use its tools!!  
We need the following two **assumptions**:

Access to a **regression oracle** SqAlg. At each round the regressor is trained using the feedback of the first  $t - 1$  rounds and outputs a regressor  $f_t$ .

**Realizability assumption:** there exists a regressor  $f^* \in \mathcal{F}$  such that for all  $t$ ,  
 $f^*(x, a) = \mathbb{E}[r_t(a) | x_t = x]$ .

---

**Algorithm 1** SquareCB

---

1: **parameters:**

Learning rate  $\gamma > 0$ , exploration parameter  $\mu > 0$ .

Online regression oracle **SqAlg**.

2: **for**  $t = 1, \dots, T$  **do**

3:   Receive context  $x_t$ .

    // Compute oracle's predictions  $\widehat{y}_t(x, a) := \text{SqAlg}_t(x, a; (z_1, y_1), \dots, (z_{t-1}, y_{t-1}))$ .

4:   For each action  $a \in \mathcal{A}$ , compute  $\widehat{y}_{t,a} := \widehat{y}_t(x_t, a)$ .

5:   Let  $b_t = \arg \min_{a \in \mathcal{A}} \widehat{y}_{t,a}$ .

6:   For each  $a \neq b_t$ , define  $p_{t,a} = \frac{1}{\mu + \gamma(\widehat{y}_{t,a} - \widehat{y}_{t,b_t})}$ , and let  $p_{t,b_t} = 1 - \sum_{a \neq b_t} p_{t,a}$ .

7:   Sample  $a_t \sim p_t$  and observe loss  $\ell_t(a_t)$ .

8:   Update **SqAlg** with example  $((x_t, a_t), \ell_t(a_t))$ .

---

(pseudo-code is for minimization. It works analogously replacing losses with rewards)

## Regret guarantees

The regret of the algorithm will depend on the regret of the sequence  $\{f_t\}$ . We assume that:

$$\sum_{t \in [T]} (\hat{y}_{t,a_t} - f^*(x_t, a_t))^2 \leq \text{Reg}_{Sq}$$

Then, SquareCB guarantees regret:

$$R_T = O\left(\sqrt{K T \text{Reg}_{Sq}}\right)$$

# References

Yasin Abbasi-Yadkori, Dávid Pál, and Csaba Szepesvári. Improved algorithms for linear stochastic bandits. *Advances in neural information processing systems*, 24, 2011.

Dylan Foster and Alexander Rakhlin. Beyond ucb: Optimal and efficient contextual bandits with regression oracles. In *International Conference on Machine Learning*, pages 3199–3210. PMLR, 2020.